

Learn RISC-V CPU Implementation and **BSV**

(**BSV**: a High-Level Hardware Design Language)

Rishiyur S. Nikhil

L2: Overview of the RISC-V ISA



Reminders

Please git clone: https://github.com/rsnikhil/Learn_Bluespec_and_RISCV_Design
(git pull for latest version). Repository structure:

```
./Book_BLang_RISCV.pdf
  Slides/
    Slides_01_Intro.pdf
    Slides_02_ISA.pdf
    ...
  Exercises/
    Ex-03-A-Hello-World/
    Ex-03-B-Top-and-DUT/
    ...
  Code/
    src_Top/
    src_Drum/
    src_Fife/
    src_Common/
    ...
  Doc/Installing_bsc_Verilator_etc.{adoc,html}
```

- Slides and Exercise are numbered in sync with book Chapter numbers.
- For Exercises, please see Appendix E of the book. Some (not all) exercises have associated code in the Exercises/ directory.

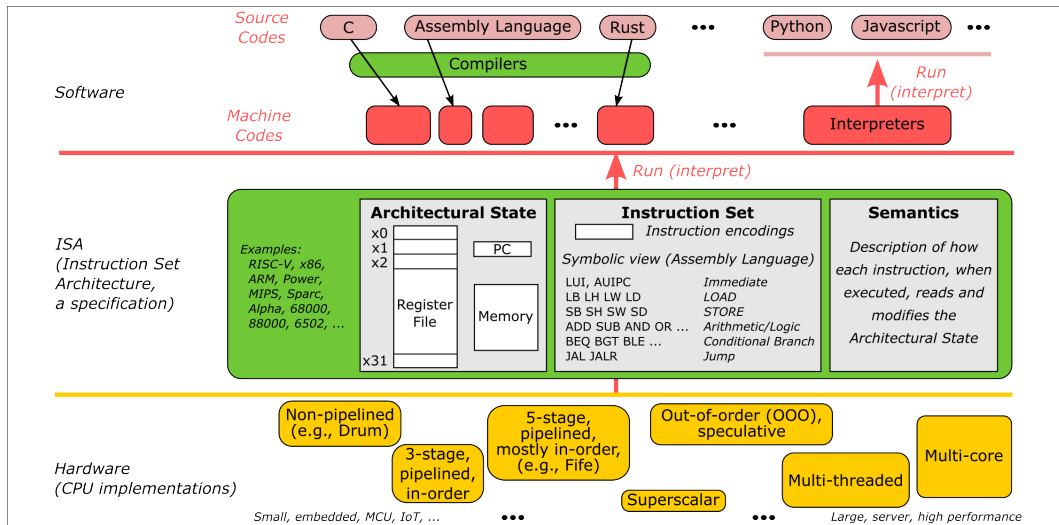
To compile and run the code for exercises, Drum and Fife, please make sure you have installed:

- bsc compiler (see <https://github.com/B-Lang-org/bsc>)
- Verilator compiler (see <https://www.verilator.org/>)

Chapter Roadmap



What is an ISA?



Architectural State

The “architectural state” is the state that is visible to instructions. For RV32I, these are:

- The PC (program counter)
- 32 General-Purpose Registers (GPRs, the “register file”)
- Memory (byte-addressed)

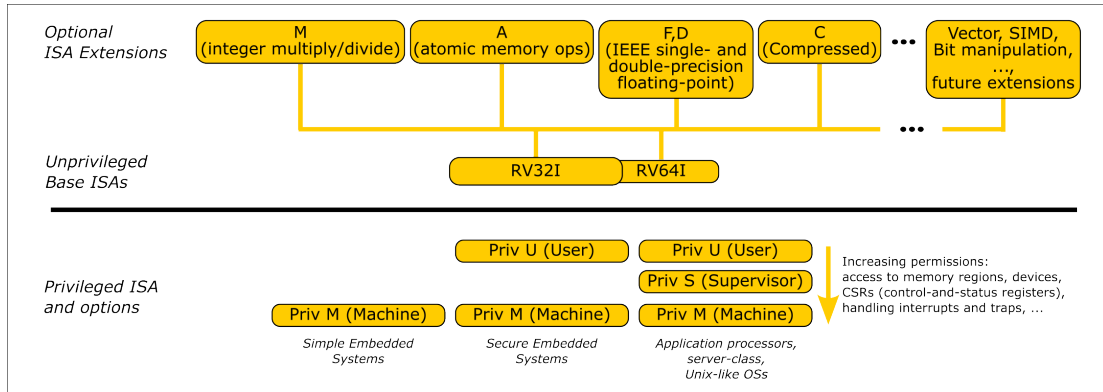
(More architectural state is defined for RV64I, and for most extensions A, F, D, Vector, ...)

The architectural state *does not include* other registers, buffers, FIFOs, memories that may be present in an implementation (they are not visible to instructions).

As such, the architectural state is present in every RISC-V implementation, from tiny CPUs for IoT devices to massive warehouse-scale servers.

Compilers only care about/know about architectural state.

Modularity of the RISC-V ISA



RISC-V Instruction Encodings

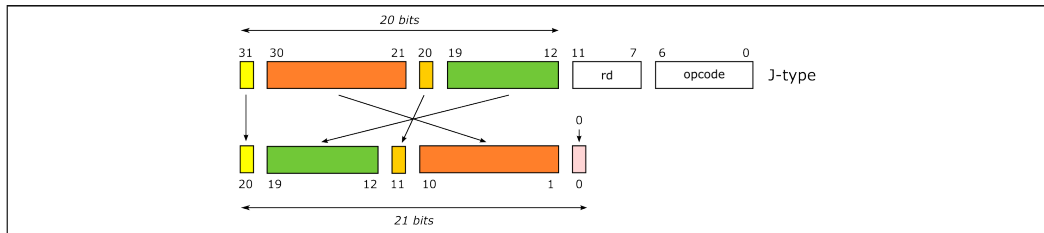
From the RISC-V specification documents:

130

Volume I: RISC-V Unprivileged ISA V20191213

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
funct7				rs2		rs1		funct3		rd		opcode		R-type
imm[11:0]						rs1		funct3		rd		opcode		I-type
imm[11:5]				rs2		rs1		funct3		imm[4:0]		opcode		S-type
imm[12 10:5]				rs2		rs1		funct3		imm[4:1 11]		opcode		B-type
imm[31:12]										rd		opcode		U-type
imm[20 10:1 11 19:12]										rd		opcode		J-type

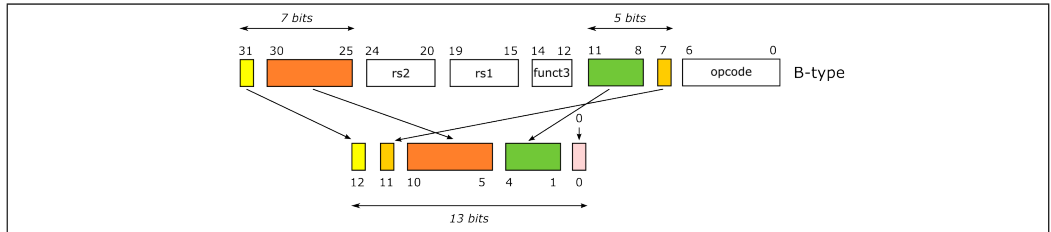
RISC-V Instruction Encodings; J-type immediates



For JAL instruction

RISC-V Instruction Encodings; B-type immediates

For BRANCH set of instructions (BEQ, BNE, BLT, BGE, BLTU, BGEU)



RV32I Instructions

RV32I Base Instruction Set

imm[31:12]				rd		0110111	LUI	↑ "load immediate"-kind: to load constant values into a register		
imm[31:12]				rd		0010111	AUIPC			
imm[20:10:11]19:12				rd		1101111	JAL	↑ "jump-and-link"-kind: subroutine calls and returns; distant jumps		
imm[11:0]				rd		1100111	JALR			
imm[12:10:5]		rs2	rs1	000	imm[4:1:11]		1100011	BEQ	↑ "conditional branch"-kind: test and possibly jump up to ~0x1000 distance	
imm[12:10:5]		rs2	rs1	001	imm[4:1:11]		1100011	BNE		
imm[12:10:5]		rs2	rs1	100	imm[4:1:11]		1100011	BLT		
imm[12:10:5]		rs2	rs1	101	imm[4:1:11]		1100011	BGE		
imm[12:10:5]		rs2	rs1	110	imm[4:1:11]		1100011	BLTU		
imm[12:10:5]		rs2	rs1	111	imm[4:1:11]		1100011	BGEU	↑ "load data from memory into register" (rs1 and imm specify address)	
imm[11:0]			rs1	000	rd		0000011	LB		
imm[11:0]			rs1	001	rd		0000011	LH		
imm[11:0]			rs1	010	rd		0000011	LW		
imm[11:0]			rs1	100	rd		0000011	LBU		
imm[11:0]			rs1	101	rd		0000011	LHU	↑ "store data from register rs2 to memory" (rs1 and imm specify address)	
imm[11:5]		rs2	rs1	000	imm[4:0]		0100011	SB		
imm[11:5]		rs2	rs1	001	imm[4:0]		0100011	SH		
imm[11:5]		rs2	rs1	010	imm[4:0]		0100011	SW		
imm[11:0]			rs1	000	rd		0010011	ADDI		↑ "integer arithmetic operations (register-immediate)"
imm[11:0]			rs1	010	rd		0010011	SLTI		
imm[11:0]			rs1	011	rd		0010011	SLTIU		
imm[11:0]			rs1	100	rd		0010011	XORI		
imm[11:0]			rs1	110	rd		0010011	ORI		
imm[11:0]			rs1	111	rd		0010011	ANDI	↑ "integer arithmetic operations (register-register)"	
0000000		shamt	rs1	001	rd		0010011	SLLI		
0000000		shamt	rs1	101	rd		0010011	SRLI		
0100000		shamt	rs1	101	rd		0010011	SRAI		
0000000		rs2	rs1	000	rd		0110011	ADD		↑ "integer arithmetic operations (register-register)"
0100000		rs2	rs1	000	rd		0110011	SUB		
0000000		rs2	rs1	001	rd		0110011	SLL		
0000000		rs2	rs1	010	rd		0110011	SLT		
0000000		rs2	rs1	011	rd		0110011	SLTU		
0000000		rs2	rs1	100	rd		0110011	XOR		
0000000		rs2	rs1	101	rd		0110011	SRL		
0100000		rs2	rs1	101	rd		0110011	SRA		
0000000		rs2	rs1	110	rd		0110011	OR		
0000000		rs2	rs1	111	rd		0110011	AND		
fm	pred	succ	rs1	000	rd		0001111	FENCE	↑ "system" operations (ignore FENCE for now)	
000000000000			00000	000	00000		1110011	ECALL		
000000000001			00000	000	00000		1110011	EBREAK		


Example specifications

Excerpt from text of ISA specification document for LUI and AUIPC instructions

Volume I: RISC-V Unprivileged ISA V20191213

19

SRLI is a logical right shift (zeros are shifted into the upper bits); and SRAI is an arithmetic right shift (the original sign bit is copied into the vacated upper bits).



LUI (load upper immediate) is used to build 32-bit constants and uses the U-type format. LUI places the U-immediate value in the top 20 bits of the destination register *rd*, filling in the lowest 12 bits with zeros.

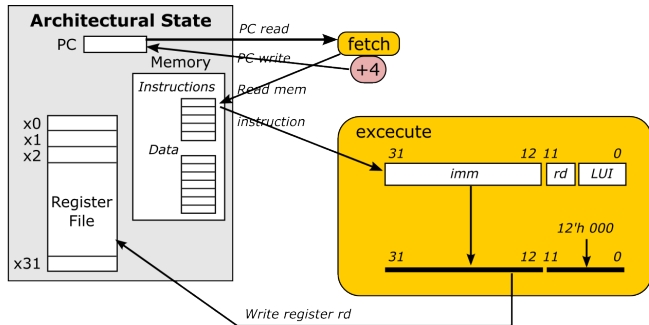
AUIPC (add upper immediate to *pc*) is used to build *pc*-relative addresses and uses the U-type format. AUIPC forms a 32-bit offset from the 20-bit U-immediate, filling in the lowest 12 bits with zeros, adds this offset to the address of the AUIPC instruction, then places the result in register *rd*.

Execution semantics: example

Execution semantics of LUI instruction

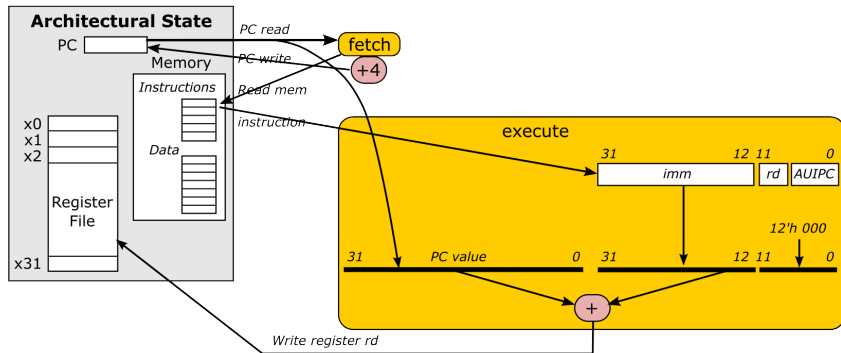
CPU operation

loop forever:
instr = fetch (PC)
execute (instr)




Execution semantics: example

Execution semantics of AUIPC instruction




Execution semantics: example

Execution semantics of JAL and JALR instructions (unconditional jumps)

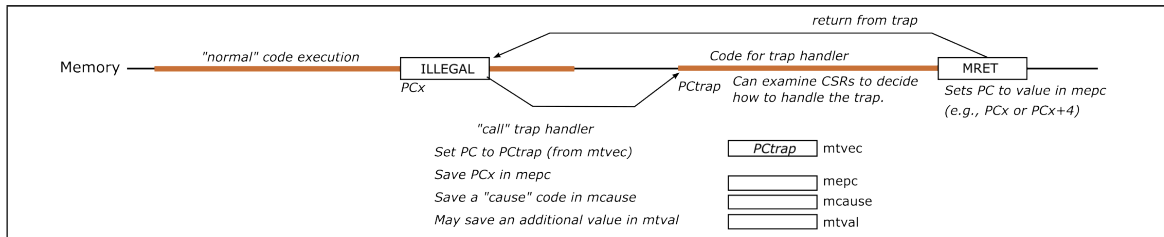


Control and Status Registers (CSRs) for Trap-Handling



These CSRs are implicitly read and written when taking and returning from a trap.
These CSRs can be explicitly read and written by CSRRxx instructions.

Trap and Trap-return flow



There are many possible causes for exceptions and traps; the illustration is for an illegal instruction.

Exception causes

On a trap, a cause-code is written into the `mcause` CSR:

Exception-Cause code	Description
0	Instruction address misaligned
1	Instruction access fault
2	Illegal instruction
3	Breakpoint
4	Load address misaligned
5	Load access fault
6	Store/AMO address misaligned
7	Store/AMO access fault
...	...
11	Environment call M-mode
...	...

CSRRxx Instructions

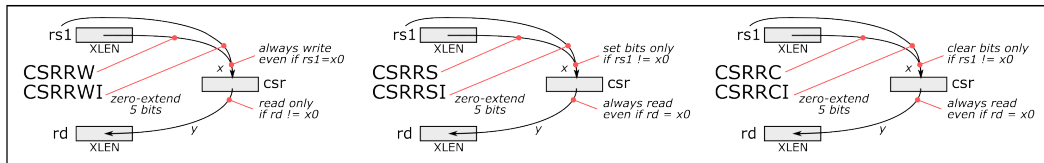
For reading and writing CSRs from RISC-V code
(from the RISC-V ISA specifications document)

9.1 CSR Instructions

All CSR instructions atomically read-modify-write a single CSR, whose CSR specifier is encoded in the 12-bit *csr* field of the instruction held in bits 31–20. The immediate forms use a 5-bit zero-extended immediate encoded in the *rs1* field.

31	20 19	15 14	12 11	7 6	0
csr		rs1	funct3	rd	opcode
12		5	3	5	7
source/dest		source	CSRRW	dest	SYSTEM
source/dest		source	CSRRS	dest	SYSTEM
source/dest		source	CSRRC	dest	SYSTEM
source/dest		uimm[4:0]	CSRRWI	dest	SYSTEM
source/dest		uimm[4:0]	CSRRSI	dest	SYSTEM
source/dest		uimm[4:0]	CSRRCI	dest	SYSTEM

CSRRxx Instruction Semantics



- They all move potentially data into a CSR, and from a CSR to a GPR
- The non “I” variants take input from GPR[rs1] (unless rs1 is zero)
- The “I” variants use rs1 itself as input (unless rs1 is zero)

RV64I instructions

From the RISC-V ISA specifications document

Volume I: RISC-V Unprivileged ISA V20191213

131

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
funct7				rs2		rs1		funct3		rd		opcode		R-type
imm[11:0]						rs1		funct3		rd		opcode		I-type
imm[11:5]				rs2		rs1		funct3		imm[4:0]		opcode		S-type

RV64I Base Instruction Set (in addition to RV32I)

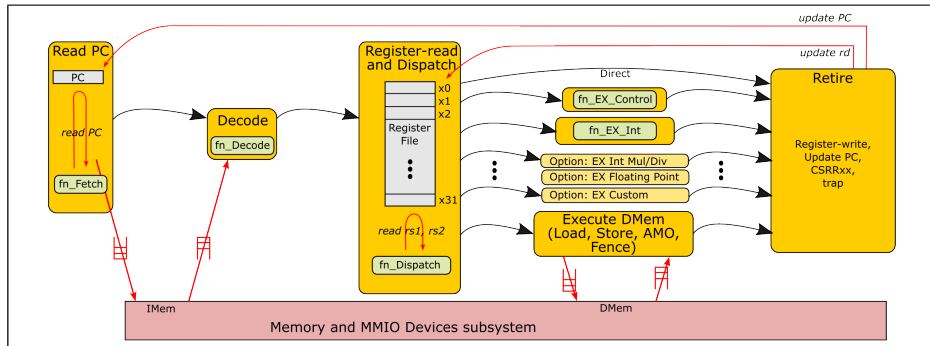
imm[11:0]			rs1	110	rd	0000011	LWU
imm[11:0]			rs1	011	rd	0000011	LD
imm[11:5]		rs2	rs1	011	imm[4:0]	0100011	SD
000000		shamt	rs1	001	rd	0010011	SLLI
000000		shamt	rs1	101	rd	0010011	SRLI
010000		shamt	rs1	101	rd	0010011	SRAI
imm[11:0]			rs1	000	rd	0011011	ADDIW
0000000		shamt	rs1	001	rd	0011011	SLLIW
0000000		shamt	rs1	101	rd	0011011	SRLIW
0100000		shamt	rs1	101	rd	0011011	SRAIW
0000000		rs2	rs1	000	rd	0111011	ADDW
0100000		rs2	rs1	000	rd	0111011	SUBW
0000000		rs2	rs1	001	rd	0111011	SLLW
0000000		rs2	rs1	101	rd	0111011	SRLW
0100000		rs2	rs1	101	rd	0111011	SRAW

RV64I instructions

Architectural state: In RV64, the 32 GPRs (General Purpose Registers) and the PC are each 64-bits wide.

- Most of the RV32I instructions have identical RV64I counterparts (here they operate on 64-bit values).
- A few instructions (SLLI, SRLI, SRAI) are slightly different (allowing 6 bits instead of 5 for the shift amount).
- A few instructions are new, to operate on 32-bit values in the 64-bit registers.
 - LWU to move a 32-bit value from memory into a 64-bit register
 - LD to load a 64-bit value from memory to a 64-bit register
 - SD to store a 64-bit value to memory from 64-bit register
 - ADDIW, SLLIW, ... SRAW to operate on 32-bits of 64-bit registers

Abstract algorithm for interpreting an ISA



This is “abstract” in the sense that it just describes necessary functionality. Different implementations will make choices as to whether or not these functions are pipelined; if pipelined, how many stages; whether or not there are concurrent pipelines; *etc.*

End