Learn RISC-V CPU Implementation and BSV
(BSV: a High-Level Hardware Design Language)

Rishiyur S. Nikhil

L5: **BSV** Structs; Memory requests and responses

bluespec ⢈

Reminders

Please git clone: https://github.com/rsnikhil/Learn_Bluespec_and_RISCV_Design
(git pull for latest version). Repository structure:

```
./Book_BLang_RISCV.pdf
  Slides/
      Slides_01_Intro.pdf
      Slides_02_ISA.pdf
      ...
  Exercises/
      Ex-03-A-Hello-World/
      Ex-03-B-Top-and-DUT/
      ...
  Code/
      src_Top/
      src_Drum/
      src_Fife/
      src_Common/
      ...
  Doc/Installing_bsc_Verilator_etc.{adoc,html}
```

- Slides and Exercise are numbered in sync with book Chapter numbers.
- For Exercises, please see Appendix E of the book. Some (not all) exercises have associated code in the Exercises/ directory.
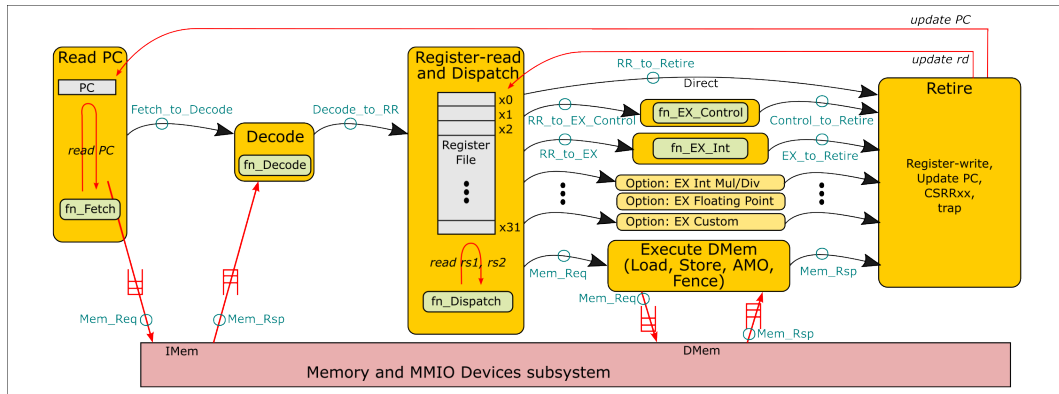
To compile and run the code for exercises, Drum and Fife, please make sure you have installed:

- *bsc* compiler (see https://github.com/B-Lang-org/bsc)
- Verilator compiler (see https://www.verilator.org/)

Chapter Roadmap

BSV topics

Packages	Ch 3	imports, exports
Combinational circuits	Ch 4	Bit-vectors, slicing, operators, integers, identifiers, booleans, functions
		simple linear FSMs
		conditionals (muxes)
Structs and Tuples	Ch 5	
Modules and Interfaces	Ch 7	Registers, Register Files, FIFOs
Sequential FSMs	Ch 9	sequences, conditionals, while-loops
Verification	Ch 11	sequences, conditionals, while-loops
BSV Rules	Ch 13	Rule conditions, atomicity, mapping rules to clocks
Concurrent Registers	Ch 17	CRegs, Pipeline and Bypass FIFOs
Suggested further study	Ch 19	Higher-order functions, Polymorphism, typeclasses, Tagged unions, pattern matching, Multiple clock domains, Bluecheck

RISC-V examples and topics

Ch 2	RISC-V ISA	
	Recognize a legal opcode	
	Simple testbench	
	Classify an instruction	
	Memory requests/responses	
Ch 6	RISC-V core combinational functions	fn_Fetch, fn_Decode, fn_Dispatch, fn_Exec_Int, fn_Exec_Control
Ch 8	GPRs and CSRs	
Ch 10	Drum CPU	Trap-handling
Ch 12	CPU verification	
Ch 14	Drum CPU with Rules instead of StmtFSM	
Ch 15-16	Fife CPU	PC prediction, Register R/W hazards, Concurrent execution pipelines, Speculative memory access, In-order retire and trap-handling
Ch 18	Fife and Drum Optimization	Eliminate pipeline bubbles, Fuse rules (reduce cycles), Reduce state, improve PC prediction, Bypassing, Caches
Ch 20	Suggested further study	RV64I, M, A, F,D, C extensions, M,S,U privilege levels, Interrupts, Virtual memory, PMPs and protection, Memory systems, Linux-capability

Flow of information between stages in Drum and Fife

The green annotations indicate the type of information flowing on each arrow.
Each of these is a "struct" type (also known as a "record"): a grouping of *fields* of heterogeneous types.

Memory requests

```
────────────── src_Common/Mem_Req_Rsp.bsv: line 44 ... ──────────────
typedef struct {Mem_Req_Type  req_type;
                Mem_Req_Size  size;
                Bit #(64)     addr;
                Bit #(64)     data;    // CPU => mem data
                ...
} Mem_Req
deriving (Bits, FShow);
```

We will discuss Mem_Req_Type and Mem_Req_Size in some following slides (they are just small bit-width scalars).
We will discuss the choice of 64 bits for addr and data in some following slides.

The data field is only relevant when communicating data from the CPU to memory (STORE and AMO instructions).

When communicating 1, 2 and 4 bytes, these are in the least-significant bytes of the data field.

Note, we do not say "deriving (Eq)" because we have no occasion to compare two entire Memory Requests for
equality/inequality.

BSV: struct expressions to construct struct values

```
Mem_Req x = Mem_Req {req_type: funct5_LOAD,
                     size:     MEM_SIZE_4B,
                     addr:     'h_8000_0000,
                     data:     ?,
                     ...    };
```

The right-hand side is a "struct expression" whose value is a struct value.

If a field is left undefined, *bsc* will warn while compiling.

Some shorthands: "let" and don't-care values:

```
let  x   = Mem_Req {req_type: funct5_LOAD,
                    size:     MEM_SIZE_4B,
                    addr:     'h_8000_0000,
                    data:     ?,
                    ...    };
```

BSV: Accessing and updating struct fields

These notations are standard in many programming languages.

Accessing fields:

```
x.req_type
x.size
```

Updating fields:

```
x.req_type = MEM_REQ_STORE;
x.data     = ... new value ... ;
```

BSV: Bit-representation of struct values

- Because we said "deriving (Bits)", *bsc* will automatically pick a hardware representation of this struct as a bit-vector, by simply concatenating the bit-representations of the fields. So, the size of the bit-vector for a struct value is the sum of the sizes of the bit-vectors for the fields.

- If we wanted a different, custom representation, we omit "deriving (Bits)", and there is a way ("Typeclass Instances") to specify exactly what we want.

BSV: printing/logging struct values (for debugging)

We can print a struct value directly, *e.g.*,

```
Mem_Req mem_req;
...
$display ("mem_req is: ", mem_req);
```

This will just print the hexadecimal notation for the full bit-vector representing the struct. This can be difficult to read:

- Some structs are large (hundreds of bits!)
- Field boundaries may not align with hexadecimal bit boundaries (every 4 bits), and so correlating the hex digits to the fields can be tedious.

BSV: printing/logging struct values (for debugging)

Because we said "deriving (FShow)", *bsc* will automatically define an "fshow() function for this struct, that will print each field separately.

```
Mem_Req mem_req;
...
$display ("mem_req is: ", fshow (mem_req));
```

If we wanted way to print the struct in a custom format, we omit "deriving (FShow)", and there is a way ("Typeclass Instances") to define fshow() to print what we want.

Memory requests: `Mem_Req_Type`

We could define the memory request-type as an enum:

```
typedef enum { MEM_REQ_LOAD, MEM_REQ_STORE} Mem_Req_Type
deriving (Bits, FShow, Eq);
```

However, looking ahead to future support of the "A" extension (Atomic Memory Ops (AMOs), see Unprivileged ISA Spec p.132), we observe that the ISA defines a 5-bit code in the instruction for each AMO op. If we use these 5 bits as-is, it will simplify decoding. So, we use 5 bits for the memory request-type:

—————————————————— src_Common/Mem_Req_Rsp.bsv: line 16 ... ——————————————————
```
typedef Bit #(5) Mem_Req_Type;
```

and we pick two 5-bit codes that are unused by the AMO ops for LOAD and STORE:

—————————————————— src_Common/Instr_Bits.bsv: line 226 ... ——————————————————
```
Bit #(5) funct5_LOAD    = 5'b_11110;
Bit #(5) funct5_STORE   = 5'b_11111;
Bit #(5) funct5_FENCE   = 5'b_11101;
```

Memory requests: Mem_Req_Size

```
─────────────── src_Common/Mem_Req_Rsp.bsv: line 41 ... ───────────────
typedef enum {MEM_1B, MEM_2B, MEM_4B, MEM_8B} Mem_Req_Size
deriving (Bits, FShow, Eq);
```

Why MEM_8B? (RV32I can only LOAD/STORE bytes (1 byte), halfwords (2 bytes) and words (4 bytes).)

This is with an eye towards future extension of our implementation:

- If we support the D ISA extension (double-precision floating point), we'll need to be able to load/store doublewords (8 bytes). This could be in RV32 or RV64.
- If we support RV64I, we'll need to be able to load/store doublewords (8 bytes).
- Even though RV32I and RV64I instructions are ≤ 32-bits wide, Fetch may choose to fetch more bits on each memory access, effectively "pre-fetching" subsequent instruction and reducing the number of memory accesses.

Why is the addr field 64 bits instead of 32? This is again with an eye towards future extension of our implementation:

- If we support RV64I, addresses will be 64 bits.
- Even in RV32, physical memory/devices can allocated be at addresses larger than the 2^{32} address space.

Please see Appendix E, Section Ex-05-A-Structs.

Memory responses

Memory responses may report an exception (misaligned access, non-existent memory, ...).

```
──────────────────── src_Common/Mem_Req_Rsp.bsv: line 58 ... ────────────────────
typedef enum {MEM_RSP_OK,
              MEM_RSP_MISALIGNED,
              MEM_RSP_ERR,
              ...
} Mem_Rsp_Type
deriving (Bits, FShow, Eq);
```

```
──────────────────── src_Common/Mem_Req_Rsp.bsv: line 67 ... ────────────────────
typedef struct {Mem_Rsp_Type  rsp_type;
                Bit #(64)     data;      // mem => CPU data
                ...
} Mem_Rsp
deriving (Bits, FShow);
```

The data field is only relevant when communicating data from memory to the CPU (LOAD and LR instructions). When communicating 1, 2 and 4 bytes, these are in the least-significant bytes of the data field.

Please see Appendix E, Section Ex-05-B-Mem-Req-Rsp.

BSV: Tuples: pre-defined immutable structs with special notation

Constructing a 2-tuple value: Example:

```
                              ─── from src_Common/CSRs.bsv ───
   function ActionValue #(Tuple2 #(Bool, Bit #(XLEN)))
           fav_csr_read (Bit #(12) csr_addr);
     ...
        return tuple2 (exception, y);
   endfunction
```

Accessing struct components using predefined functions tpl_*j*:

```
   let xy <- fav_csr_read (...);
   let exc = tpl_1 (xy);    // exc has type: Bool
   let v   = tpl_2 (xy);    // v   has type: Bit #(XLEN)
```

Accessing struct components using pattern-matching:

```
   match { .exc, .v } <- fav_csr_read (csr_addr);
```

"Harvard" Architecture: Separating Instruction and Data Memory

Since the very days of computers, most computers have separate channels to memory:

- "IMem": for Fetch to read from instruction memory
- "DMem": for LOAD/STORE instructions to read/write data memory

Typically, IMem and DMem can be accessed concurrently.

To facilitate this concurrency, programs typically do not modify their own instructions with LOAD/STORE instructions.
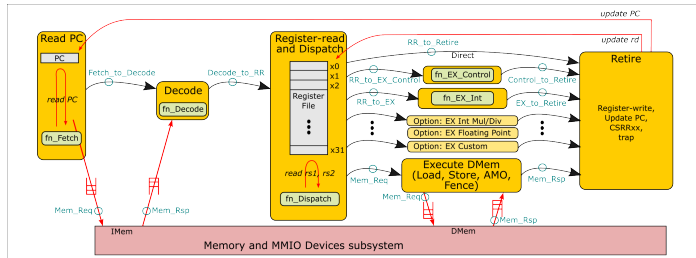
This organization is sometimes loosely called a "Harvard Architecture".
See: https://en.wikipedia.org/wiki/Harvard_architecture.

End