Learn RISC-V CPU Implementation and **BSV**
(**BSV**: a High-Level Hardware Design Language)

Rishiyur S. Nikhil

L17: **BSV**: Tighter Rule scheduling with CRegs

bluespec ⚙

Reminders

Please git clone: `https://github.com/rsnikhil/Learn_Bluespec_and_RISCV_Design`
(git pull for latest version). Repsitory structure:

```
./Book_BLang_RISCV.pdf
  Slides/
      Slides_01_Intro.pdf
      Slides_02_ISA.pdf
      ...
  Exercises/
      Ex-03-A-Hello-World/
      Ex-03-B-Top-and-DUT/
      ...
  Code/
      src_Top/
      src_Drum/
      src_Fife/
      src_Common/
      ...
  Doc/Installing_bsc_Verilator_etc.{adoc,html}
```

- Slides and Exercise are numbered in sync with book Chapter numbers.
- For Exercises, please see Appendix E of the book. Some (not all) exercises have associated code in the Exercises/ directory.

To compile and run the code for exercises, Drum and Fife, please make sure you have installed:

- *bsc* compiler (see `https://github.com/B-Lang-org/bsc`)
- Verilator compiler (see `https://www.verilator.org/`)

Chapter Roadmap

BSV topics

Packages — **Ch 3** — imports, exports

Combinational circuits — **Ch 4** — Bit-vectors, slicing, operators, integers, identifiers, booleans, functions

simple linear FSMs
conditionals (muxes)

Structs and Tuples — **Ch 5**

Modules and Interfaces — **Ch 7** — Registers, Register Files, FIFOs

Sequential FSMs — **Ch 9** — sequences, conditionals, while-loops

Verification — **Ch 11** — sequences, conditionals, while-loops

BSV Rules — **Ch 13** — Rule conditions, atomicity, mapping rules to clocks

Concurrent Registers — **Ch 17** — CRegs, Pipeline and Bypass FIFOs

Suggested further study — **Ch 19** — Higher-order functions, Polymorphism, typeclasses, Tagged unions, pattern matching, Multiple clock domains, Bluecheck

RISC-V examples and topics

Ch 2 — RISC-V ISA

Recognize a legal opcode

Simple testbench
Classify an instruction

Memory requests/responses

Ch 6 — RISC-V core combinational functions — fn_Fetch, fn_Decode, fn_Dispatch, fn_Exec_Int, fn_Exec_Control

Ch 8 — GPRs and CSRs

Ch 10 — Drum CPU — Trap-handling

Ch 12 — CPU verification

Ch 14 — Drum CPU with Rules instead of StmtFSM

Ch 15-16 — Fife CPU — PC prediction, Register R/W hazards, Concurrent execution pipelines, Speculative memory access, In-order retire and trap-handling

Ch 18 — Fife and Drum Optimization — Eliminate pipeline bubbles, Fuse rules (reduce cycles), Reduce state, improve PC prediction, Bypassing, Caches

Ch 20 — Suggested further study — RV64I, M, A, F,D, C extensions, M,S,U privilege levels, Interrupts, Virtual memory, PMPs and protection, Memory systems, Linux-capability

Table of Contents

BSV: Tighter Rule Scheduling: Motivation

Tighter Rule Scheduling: Motivation

|	**BSV**	*Analogy: RISC-V ISA*
(A)	Rule semantics are abstract: one rule at a time	*RISC-V ISA semantics are abstract: one instruction at a time*
(B)	When the *bsc* compiler maps a set of rules to execute simultaneously (in the same clock), all their methods occur at the same instant (clock edge). All "read-values" are from the previous clock edge; all "write-values" (`Action` and `ActionValue` results) are visible only at the next clock edge.	*RISC-V instructions can be executed in pipelines, in parallel (superscalar), out-of-order, ... The implementation must ensure that the order in which they retire, and the order in which they read and write registers and memory, is consistent with (A).*
	bsc must ensure that the *ordering* of methods in (B) is consistent with (A).	*In particular, an instruction may be stalled to avoid violating (A).*
	In particular, *bsc* may introduce (combinational) control circuits to *stall* a rule to avoid violating (A).	

Tighter Rule Scheduling: Motivational Example: Up-Down Counter (1/2)

Consider this scenario, where a Producer streams packets over a long connection to a Consumer:

To avoid over-running the receive buffer, The producer:

- initializes the counter to the available space in the buffer;
- for every packet sent, decrements the counter (`decr` method);
- stalls (doesn't send) if the counter value (`val` method) is zero (no space available).

As the receiver consumes packets from the buffer, it sends acknowledgements back to indicate the amount of space freed.

In the transmitter, acknowledgements are incremented back into the counter (`incr` method), allowing transmission to continue.

```
interface Up_Down_Counter_IFC;
   method Action   init (Bit #(4) init_val);
   method Bit #(4) val;
   method Action   decr;
   method Action   incr;
endinterface
```

Here is a possible implementation of the "Up-Down" Counter:

```
module mkUp_Down_Counter_I (Up_Down_Counter_IFC);
   // STATE
   Reg #(Bit #(4)) rg_counter <- mkReg (15);

   // --------------------------------
   // INTERFACE

   method Action init (Bit #(4) init_val);
      rg_counter <= init_val;
   endmethod

   method Bit #(4) val;
      return rg_counter;
   endmethod

   method Action decr () if (rg_counter != 0);
      rg_counter <= rg_counter - 1;
   endmethod

   method Action incr () if (rg_counter != 15);
      rg_counter <= rg_counter + 1;
   endmethod
endmodule
```

Analysis: Rules invoking .incr() and .decr() cannot execute on the same clock edge.

- (A) In the abstract rule semantics, either .incr() precedes .decr() or *vice versa*.

  In either case, the latter rule observes the update from the previous rule.

- (B) If they executed on the same clock, neither rule sees the update from the other rule. This is inconsistent with (A).

Thus, this implementation cannot send a packet (decr) and register an ack (incr) in the same clock, even though the two streams are asynchronous and concurrent.

On the average, we can only send (and register ack) on every alternate clock.

Solution:
Use a *Concurrent Register* (CReg) for rg_counter.

BSV: Concurrent Registers (CRegs)

CRegs

A Concurrent Register, or CReg, is a module provided by the *bsc* library. Its interface is an *array* of Reg#(t) interfaces:

```
module mkCReg #(parameter Integer n,
                parameter a_type resetval)
              (Reg#(a_type) ifc[])
```

It is instantiated similar to this example:
(This example: register content type: Bit#(4); initial value: 15; two Reg interfaces).

```
    // parameter n is 2, resetval is 15
    Array #(Reg #(Bit #(4))) crg_counter <- mkCReg (2, 15);
```

Hardware intuition for a CReg
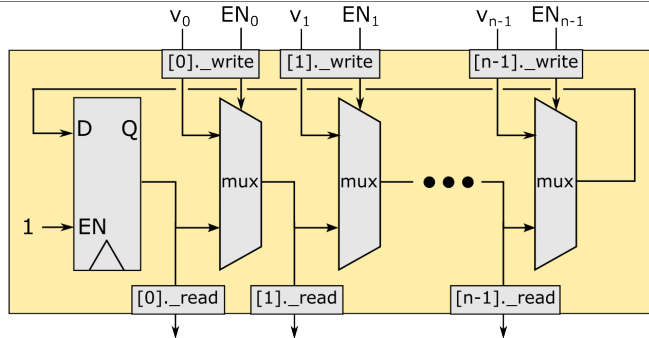
This has an array of n Reg interfaces, indexed from 0 to n-1.

The j'th interface has [j]._read() and [j]._write() methods.

CReg ordering properties

- All the methods can be invoked in the same clock.
- A read at the j'th register interface, *i.e.*, $x[j].\_read$ returns the latest of:
  - the value v_{j-1}, if $x[j-1].\_write(v_{j-1})$ is being invoked;
  - else the value v_{j-2}, if $x[j-2].\_write(v_{j-2})$ is being invoked;
  - ...
  - else the value v_1, if $x[1].\_write(v_1)$ is being invoked;
  - else the value v_0, if $x[0].\_write(v_0)$ is being invoked;
  - else the value in the register.
- The register value is updated with the latest of:
  - the value v_{n-1}, if $x[n-1].\_write(v_{n-1})$ is being invoked;
  - else the value v_{n-2}, if $x[n-2].\_write(v_{n-2})$ is being invoked;
  - ...
  - else the value v_1, if $x[1].\_write(v_1)$ is being invoked;
  - else the value v_0, if $x[0].\_write(v_0)$ is being invoked;
  - else the current value in the register.

This ordering corresponds exactly to the left-to-right and top-to-bottom ordering of the methods in the hardware-intuition diagram (Slide 11).

Up-down counter with a CReg

```
module mkUp_Down_Counter_I (Up_Down_Counter_IFC);
   // STATE
   Array #(Reg #(Bit #(4))) crg_counter <- mkCReg (3,15);

   // -------------------------------
   // INTERFACE
   method Action init (Bit #(4) init_val);
      crg_counter [2] <= init_val
   endmethod

   method Bit #(4) val;
      return rg_counter [0];
   endmethod

   method Action decr if (rg_counter != 0);
      rg_counter [1] <= rg_counter [1] - 1;
   endmethod

   method Action incr if (rg_counter != 15);
      rg_counter [0] <= rg_counter [0] + 1;
   endmethod
endmodule
```

Some questions to ponder:

- What does method val return?
  The original value in the register?
  The value after the increment?
  The value after the decrement?
  The value after the increment and decrement?

- What happens if methods incr and decr are called simultaneously? Which one happens (semantically) "first"?
  Note: incr saturates at 15, and decr saturates at 0, so the order matters!

Hint: The answers are in the choice of CReg indexes in each method.

Example: CSR `mcycle` in RISC-V

CSR `mcycle` (RISC-V CPU cycle counter) is updated by two "processes":

- (A) Standalone infinite loop incrementing `mcycle` on every cycle.
- (B) Instruction execution, when we have a `CSRRxx` instruction that writes to `mcycle`.

The RISC-V spec says that when (B) happens, it should override (A).

We can instantiate a CReg for this:

```
───────────────────────── from src_Common/CSRs.bsv ─────────────────────────
   Array #(Reg #(Bit #(64))) csr_mcycle <- mkCReg (2, 0);
```

(A) Standalone process incrementing it:

```
───────────────────────── from src_Common/CSRs.bsv ─────────────────────────
   rule rl_count_cycles;
      csr_mcycle [0] <= csr_mcycle [0] + 1;
   endrule
```

(B) CSRRxx instruction execution (overrides (A) because of higher CReg index):

```
──────────────── from src_Common/CSRs.bsv in function fav_csr_write() ────────────────
   csr_mcycle [1] <= csr_val;
```

BSV: Higher-performance FIFOs (in library `SpecialFIFOs`)

(implemented using Concurrent Registers)

A FIFO implemented with ordinary registers

Consider the following module implementing a 1-element FIFO with a FIFOF interface,
using ordinary registers (mkReg, mkRegU):

```
module mkFIFOF (FIFOF #(Bit #(32)));
   Reg #(Bit #(32)) rg_data <- mkRegU;
   Reg #(Bool)      rg_full <- mkReg (False);

   // ----------------
   // INTERFACE

   method Bool notEmpty ();
      return rg_full;
   endmethod

   method Bit #(32) first () if (rg_full);
      return rg_data;
   endmethod

   method Action deq () if (rg_full);
      rg_full <= False;
   endmethod
   ...
```

```
   ...
   method Bool notFull ();
      return (! rg_full);
   endmethod

   method Action enq (Bit #(32) x) if (! rg_full);
      rg_data <= x;
      rg_full <= True;
   endmethod

   method Action clear;
      rg_full <= False;
   endmethod
endmodule
```

Analysis of the performance of `mkFIFOF`

Consider a producer rule `rl_P` that invokes `enq`, and a consumer rule `rl_C` that invokes `first` and `deq`:

These rules cannot fire at the same instant (on the same clock) because both of them read and write register `fg_Full` (because both `enq` and `deq` read and write the register).

We can use a CReg to relax this constraint, in two different ways, which we call `mkPipelineFIFOF` and `mkBypassFIFOF`, respectively.

mkPipelineFIFOF

Instead of `mkReg` for `rg_full`, we can use `mkCReg`:

```
module mkFIFOF (FIFOF #(Bit #(32)));
module mkPipelineFIFOF (FIFOF #(Bit #(32)))
   Reg #(Bit #(t))        rg_data   <- mkRegU;
   Array #(Reg #(Bool)) crg_full <- mkCReg (3, False);

   // ----------------
   // INTERFACE

   method Bool notEmpty ();
      return crg_full [0];
   endmethod

   method Bit #(32) first () if (crg_full [0]);
      return rg_data;
   endmethod

   method Action deq () if (crg_full [0]);
      crg_full [0] <= False;
   endmethod
   ...
```

```
   ...
   method Bool notFull ();
      return (! crg_full [1]);
   endmethod

   method Action enq (Bit #(32) x) if (! crg_full [1]);
      rg_data       <= x;
      crg_full [1] <= True;
   endmethod

   method Action clear;
      crg_full [2] <= False;
   endmethod
endmodule
```

Note that the "dequeue" side methods use index [0], "enqueue" side
use methods index [1]; and "clear" uses index [2].

These choices affect the ordering semantics of the methods.

Analysis of the performance of `mkPipelineFIFOF`

Consider again a producer rule `rl_P` that invokes `enq`, and a consumer rule `rl_C` that invokes `first` and `deq`:

Because of `mkPipelineFIFOF`'s CReg indexes, in the equivalent rule-at-a-time semantics, `rl_C` fires *before* `rl_P`. Thus, even if the FIFO was full at the start of the clock, `rl_P` can still `enq` into the FIFO, *provided* `rl_C` is firing on the same clock.

Per the rule-at-a-time semantics, `rl_C` fires first, which empties the FIFO, *i.e.,* `rl_P` sees the FIFO as empty and is therefore able to `enq` into it.

This is why we call it a "PipelineFIFO": it is an ideal candidate for the FIFO between stages of a pipeline, allowing the downstreanm stage (the consumer) and the upstream stage (the producer) to fire on the same clock, advancing data in a piplined manner.

Note also that another rule invoking `clear`, too, can fire on the same clock as `rl_P` and `rl_C`. Because of its CReg index, logically it fires "last", and so will leave the FIFO in a finally empty state.

Analysis of the hardware for `mkPipelineFIFOF`

In the Verilog for `mkPipelineFIFOF`, the READY signal for `enq` incorporates the ENABLE signal of the `deq` method (because if the FIFO is full from the previous clock, in this clock `enq` can only be invoked if `deq` is also being invoked).

Thus, there is a *combinational path* backward through the FIFO (a path that involves only wires and gates and no state-element) from the `deq` method to the `enq` method.

mkBypassFIFOF

What happens if we exchange the [0] and [1] CReg indexes?

```
Array #(Reg #(Bit #(32))) crg_data <- mkCRegU (2);
Array #(Reg #(Bool))      crg_full <- mkCReg (3, False);

// ----------------
// INTERFACE

method Bool notEmpty ();
   return crg_full [1];
endmethod

method Bit #(32) first () if (crg_full [1]);
   return crg_data [1];
endmethod

method Action deq () if (crg_full [1]);
   crg_full [1] <= False;
endmethod
...
```

```
   ...
method Bool notFull ();
   return (! crg_full [0]);
endmethod

method Action enq (Bit #(32) x) if (! crg_full [0]);
   crg_data [0] <= x;
   crg_full [0] <= True;
endmethod

method Action clear;
   crg_full [0] <= False;
endmethod
endmodule
```

Now, the "enqueue" side methods use index [0] and the "dequeue" side methods use index [1] ("clear" still uses index [2]).

These choices change the ordering semantics of the methods.

Analysis of the performance of `mkBypassFIFOF`

Consider again a producer rule `rl_P` that invokes enq, and a consumer rule `rl_C` that invokes `first` and `deq`:

Because of `mkBypassFIFOF`'s CReg indexes, in the equivalent rule-at-a-time semantics, `rl_C` fires *after* `rl_P`. Thus, even if the FIFO was empty at the start of the clock, `rl_C` can still deq from the FIFO, *provided* `rl_P` is firing on the same clock.

Per the rule-at-a-time semantics, `rl_P` fires first, which fills the empty FIFO, *i.e.*, `rl_C` sees the FIFO as full and is therefore able to deq from it.

This is why we call it a "BypassFIFO". When used in a pipeline, the enqueued value can be "bypassed" straight through the FIFO to the consumer on the same clock.

As with `mkPipelineFIFOF`, another rule invoking clear, too, can fire on the same clock as `rl_P` and `rl_C`. Because of its CReg index, logically it fires "last", and so will leave the FIFO in a finally empty state.

Analysis of the hardware for `mkBypassFIFOF`

In the Verilog for `mkBypassFIFOF`, the READY signal for `first` and `deq` incorporates the ENABLE signal of the `enq` method (because if the FIFO is empty from the previous clock, in this clock `first` and `deq` can only be invoked if `enq` is also being invoked).

Similarly, for the data to be bypassed through, there has to be a combinational path from the `enq` argument to the `first` result.

Thus, there are *combinational paths* forward through the FIFO (paths that involve only wires and gates and no state-element) from the `enq` method to the `first` and `deq` methods.

hardware wires for enq

READY
ENABLE
data

READY
ENABLE

hardware wires for deq

READY
data

hardware wires for first

Summary of `mkFIFOF`, `mkPipelineFIFOF` and `mkBypassFIFOF`

	`mkFIFOF`	`mkPipelineFIFOF`	`mkBypassFIFOF`
1-tick traversal	Yes	Yes	Yes
# of buffer registers	2	1	1
Scheduling constraints	None	deq before enq	enq before deq
Through combinational circuits	None	deq $\rightarrow$ enq	enq $\rightarrow$ deq
Separate compilation of stages	No	No	No

In the next section we'll discuss an alternative: back-to-back composition of `mkBypassFIFOF` and `mkPipelineFIFOF`

BSV: Connecting pipeline stages

Back-to-back composition of BypassFIFO and PipelineFIFO (1/2)

An interesting component is a back-to-back composition of the two high-performance FIFOs we have discussed.

Consider this code fragment, illustrated below:


```
module mk... (...);
   FIFOF #(Bit #(32)) f_bypass   <- mkBypassFIFOF;
   FIFOF #(Bit #(32)) f_pipeline <- mkPipelineFIFOF;

   // Producer rule (into f_bypass's enq side)
   rule rl_P;
      ... f_bypass.enq (x);
   endrule

   // Connect f_bypass's first/deq side
   // to f_pipeline's enq side
   mkConnection (to_FIFO_O (f_bypass),
                 to_FIFO_I (f_pipeline));

   // Consumer rule (from f_pipeline's first/deq side)
   rule rl_C;
      let y = f_pipeline.first;
      f_pipeline.deq;
   endrule
endmodule
```

Back-to-back composition of BypassFIFO and PipelineFIFO (2/2)

This has some pleasant properties:

- Despite there being two FIFOs, it takes only one tick to traverse them (because of the nature of BypassFIFOs).

- There are no ordering constraints across the composed FIFOs, *i.e.,* the FIFOs do not induce any ordering constraints between the producer rule `rl_P` and the consumer rule `rl_C` (they are "conflict free"). They can fire in the same clock and can go in either logical order.

- There are no combinational paths through the pair of FIFOs! One can verify this by studying the Verilog. The *bsc* compiler also helpfully reports the absence of combinational paths.

- Enables easier separate compilation (into Verilog) of stage modules: place one of the two component FIFOs in each stage module, and use `mkConnection` in the parent module.

  (Separate compilation is possible with the other FIFOs, but is messier.)

This makes it attractive for connecting stages in a pipeline, enabling a *modular* separation of stages, where each stage can be independently verified and compiled to Verilog. We use this technique everywhere in Fife.

Summary of `mkPipelineFIFOF` and `mkBypassFIFOF` and their composition

	`mkFIFOF`	`mkPipelineFIFOF`	`mkBypassFIFOF`	Composition
1-tick traversal	Yes	Yes	Yes	Yes
# of buffer registers	2	1	1	2
Scheduling constraints	None	deq before enq	enq before deq	None
Through combinational circuits	None	deq $\rightarrow$ enq	enq $\rightarrow$ deq	None
Separate compilation of stages	No	No	No	Yes

Final comments on CRegs

- CRegs allow us to save a cycle by allowing two rules to run concurrently in the same clock where previously they had to run in separate clocks.

  Effectively, CRegs allow us safely to *fuse* the actions in two different rules into a single composite action (whenever the rule conditions allow)

- It plays a central rule in fine-tuning **BSV** designs for optimal performance, without changing the functional semantics (rule-at-a-time).

  In RISC-V designs, CRegs enable:
    - Isolation of stages (no combinational paths) while preserving pipeline speed (as discussed in Slide 27).
    - Faster PC redirection after a misprediction
    - Faster resolution of register dependencies in the scoreboard
    - Faster reorder buffers in Out-Of-Order processors
    - ... and more ...

Caveat: Before CRegs were available in **BSV***, designs used a facility called "RWires" for the same fusion optimizations. RWires are still available in* **BSV** *(see library documentation), but we recommend using CRegs in new designs. RWires muddy the otherwise clear separation of functional semantics (logical) and clocked semantics (implementation).*

End